


Ema APM Program

# **Aria** The Strategic Sourcing Analyst

Solving the Enterprise "Procurement  
Drag" to Unlock Strategic Value



# Case Prompt



Ema is a universal AI employee, with multiple personas across roles and industries to revolutionize how people work. As a PM at Ema, you will regularly encounter deep, impactful enterprise problems where one of Ema's personas can work collaboratively with existing employees to multiply their productivity. Identify one such gap in existing businesses and propose an Ema persona to provide outsized gains to the team she would be a part of.

Address the following prompts:


- Why did you choose this problem? What is the potential impact of solving it?
- How will your persona use GenAI to solve the problem? What does the end to end user journey of working with this persona look like?
- How will you test your persona and validate your hypotheses quickly?
  - Demonstrate the feasibility of your use case via sample prompts and responses on synthetic data.
- How do you measure the success of your solution?
- What are some potential pitfalls of your solution?

@Aria What problem enterprises faces with manual sourcing processes?


H

The core problem is the **"Procurement Drag"**, a term for the systemic delays and value erosion caused by slow, manual enterprise sourcing processes. This is not a simple administrative inconvenience; it is a **strategic blocker** that directly impacts **revenue and innovation**


### The Procurement Drag Cycle



### The Hidden Cost of Procurement Drag



### How current manual process looks like ?



**40-80 hrs** of **manual evaluation**, cross referencing, and scoring by high skill, expensive sourcing managers, legal teams, and IT experts.

While **manual scoring in excel** possesses risks such as human errors, subjective & inconsistent grading, and process inefficiency.

Expectations are strategy, innovation, supplier partnerships, risk management but in reality, major chunk of valuable time is spent on manual reading, data extraction, compliance checking & excel.

JTBD

Enable strategic sourcing teams to instantly understand, compare, and evaluate complex IT/SaaS vendor proposals—so they can make faster, safer, more strategic procurement decisions without the drag of manual document analysis.

## Savings per Enterprise



### Labor savings

40–80 hours/event  
× \$100/hr × 40 events =  
\$160k–\$320k saved annually



### Opportunity cost savings

Reducing procurement drag  
by 1% on \$2B annual spend =  
\$20M unlocked value

## Enterprise Budget Shift Trend

over 5 year period



\$20B

Procurement  
Software

TAM



\$5B

Autonomous  
Sourcing

SAM

- Labour and manual evaluation
- Legacy tools and maintenance
- Strategic Investments
- Procurement tech and AI

\$500M

Serviceable  
Available  
Market

SOM

# Why solve it now?

- Procurement Complexity Has Exploded
- Legacy Tools Cannot Handle Unstructured Data
- Enterprise Budgets Are Shifting to AI Efficiency
- Cost of Slow Procurement Is Now Quantified
- Market Has Shifted to Autonomous Decisioning
- Generative AI Has Reached Technical Feasibility

Global IT  
spend growth

+7.9% → \$5.43T  
in 2025 total

80%

of CPOs plan to  
deploy Generative  
AI within ~3 years

~90%

of procurement leaders are  
evaluating or using AI agents  
for sourcing

36%

of procurement orgs had meaning-  
ful Gen AI implementations

>50%

of CPOs increased S2P /  
procurement tech budgets (2025)

[Sources for all the data shared \(click to view\)](#)

# Solution?

## Aria

### The Strategic Sourcing Analyst

Autonomous agent that reads, compares, scores, and risk-flags vendor proposals so humans do the thinking & not the reading.

*a*

### Who is Aria ?


Autonomous AI employee trained in procurement best practices and contract law.

An AI Orchestrator that reasons and acts as the sourcing manager reviews and approves Aria's outputs.

### Core Function

- **Ingests:** RFP + vendor proposals
- Produces:** Compliance scores, Risk flags & Justifications with citations

*Before Aria*




#### Transactional Work

Manually read 100+ page proposals; text-match; manual scoring (40–80 hrs/event).

#### Strategic Work

Sourcing Manager: 80% grunt work → 20% negotiation

*With Aria*



#### Analysis Review

Verify Aria's scores & citations (1–2 human-hours/event).

#### Strategic Work

Sourcing Manager: 20% review → 80% negotiation & strategy



### RFP analysis: Weeks → Hours

Target: RFP → score in <24 hours (vs 45–90 days).



### Security risk detection

Time-to-risk detection: days → <15 minutes.



### Manual work ↓ Strategic bandwidth ↑

Manual review: 40–80 hrs → <10 hrs (40–70% productivity gain).

## 3 Pillars of Trust & Deployment



**Citations Mandate:** Every finding includes source snippet & page-level citation.



**The Vault:** Single-tenant VPC / on-prem option; contractual no-training clause.



**Integration-first:** Plug into SAP Ariba, Coupa, Slack/Teams.

## User Personas

**Sarah Mehta**, Strategic Sourcing Manager

Owns large-scale sourcing events for IT, software, and enterprise tools. Balances cost, compliance, and speed under tight timelines.

**Goals**

- Manage RFPs & RFQs
- Coordinate between legal, IT, and finance
- Evaluate vendor proposals and score them

**Pain Points**

- Drowning in 300-page proposals
- Manual, Excel-based scoring
- Constant follow-ups with stakeholders

**Amit Narang**, IT Security Reviewer

Reviews vendor submissions for security compliance, data privacy, and technical architecture.


**Goals**

- Check vendor InfoSec responses
- Validate compliance with internal policies (SOC2, GDPR, etc.)
- Approve or flag risky vendors

**Pain Points**

- Hard to find specific InfoSec clauses in 100-page docs
- Missed red flags buried in text
- Delayed input slows down sourcing timeline

## User Journey




Aria turns complex sourcing into a simple, conversational flow. It starts with a quick Slack command to create a workspace, where users upload their RFPs and vendor proposals.

Aria then reads, compares, and scores each proposal automatically, highlighting risks, summarising differences, and answering stakeholder questions in real time.

Once the review is done, Aria generates a complete award justification and shares it back in Slack saving weeks of manual effort and enabling fast, confident decisions.


# System Design



## Feasibility Demonstration (Demo Data)

| RFP Requirement                                     | Vendor A                                                                                                                                 | Vendor B                                                                                                |
|-----------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| Section 3.3.8: Accessibility                        | Page 42                                                                                                                                  | Page 37                                                                                                 |
| Must comply with... WCAG 2.1... and provide a VPAT. | Our platform is fully... compliant with... WCAG 2.1 Level AA... Our... VPAT is available at <a href="http://www.cera...">www.cera...</a> | Accessibility is a key part of our product roadmap... we are actively working on future enhancements... |

## @Aria Prompt & Response



The screenshot shows the Aria Prompt & Response interface, which is a chat-based system for managing procurement tasks. The interface includes a sidebar with navigation options like **Threads**, **All DMs**, **Drafts**, **Mentions & reactions**, **Saved items**, and **More**. The main chat area displays a conversation with **#procurement-team**, showing messages from **Kenny Park**, **Aria**, **Emily Anderson**, and **Aria** regarding WCAG 2.1 compliance and VPATs. The chat interface also includes a **Vendor Proposal Comparison** section with a **Last edited just now** timestamp.

## Can Aria reduce manual evaluation & scoring time by over 70%?

Sourcing teams spend 50–100 hours per RFP manually reading PDFs and scoring Excel sheets. A 70% reduction directly converts to cost savings and cycle-time gains.

## Can Aria's MC-RAG analysis match or exceed human-level accuracy in compliance checking and risk detection?

Procurement and legal teams won't adopt automation unless accuracy meets expert standards. Aria's multi-document comparative reasoning ensures contextual, evidence-linked precision.

## Will sourcing managers trust Aria's analysis, thanks to its "Citations Mandate", and feel it frees them for strategic work?

Trust drives adoption. The goal is to make Aria feel like a "strategic co-pilot," not a black box.

### Validation Plan

#### PHASE 1



**Goal:** Validate H3 (Desirability) and the "Chat-to-Canvas" journey.



**Method:** Build the front-end UI (Slackbot + Workspace) with no AI back-end. Pick 5 internal sourcing managers. A human PM will "act" as Aria, manually finding answers and posting them as the AI.



**Measure:** Qualitative feedback. " Did this felt intuitive ? ", " Did you trust the answer ? ", " Where did you get stuck ? "

#### PHASE 2



**Goal:** Validate H1 and H2 with hard data.



**Method:** Build the "Multi-Document Comparative RAG" (MC-RAG) MVP. Pick 10 closed, historical RFP events (with known outcomes and human-hours spent) from a design partner.



**Measure (Quantitative):**

- Time-to-Analysis: Aria's processing time vs. Baseline Human Hours.
- Decision Accuracy: Did Aria's top-scored vendor match the final human decision?
- Risk Detection: Did Aria find critical risks the human team missed (or vice-versa)?

## Product Initiatives

### Autonomous Proposal Analysis

- Automate requirement extraction from RFPs
- Auto-compare vendor proposals
- Generate scoring, summaries, and citations
- Reduce manual reading & spreadsheet work

### AI-guided Decision Support

- Flag risks, gaps, compliance deviations
- Provide evidence-backed recommendations
- Surface vendor insights instantly
- Generate award justification

### Intelligent Collaboration

- Centralised workspace for SMEs
- Slack-based workflow initiation
- Real-time Q&A with Aria

### Procurement Ops Efficiency

- Remove repetitive tasks
- Standardise evaluations
- Improve auditability
- Close sourcing loops faster

## KPI

- % Requirements Auto-Scored (Leading)
- RFP Parsing Time (Leading)
- Proposal Analysis Time per Event (Leading)

- Decision Match Rate (Accuracy, Lagging)
- Recommendation Acceptance Rate
- Risk Detection Recall (Quality)

- Stakeholder Response Time (Leading)
- Ad-hoc Query Handled Rate (Usage)

- Human Hours Saved per Event (Value, Lagging)
- Audit Readiness Score (Compliance)
- Events Per Month (Scale)

## North Star Metric

### Time-to-Decision Reduction (%)

Percentage reduction in the median time it takes to get from "RFP published" → "Final vendor decision" when using Aria vs. historical baseline.

 **Risk** **Impact** **Mitigations****1. Hallucination & Inaccuracy**

A GenAI hallucination is a multi-million dollar legal and financial liability. Specialised Legal RAG systems are still shown to hallucinate in 17-33% of cases (Stanford, June 2025).

**The "Citations Mandate"**

The UI/UX is the primary mitigation. Every single score, summary, and flag generated by Aria must be accompanied by a direct, 1-click, verifiable citation back to the exact source snippet in the vendor document. The human's job is to verify, not find.

**Data Confidentiality & IP Risk**

Vendors will not upload their most sensitive IP, pricing, and solution designs to a "leaky" multi-tenant model. This is an immediate non-starter for any enterprise.

**The "Vault" Architecture**

Aria must be deployed in a single-tenant, private cloud (VPC) or on-premise environment. Ema must provide a non-negotiable contractual guarantee that no customer-uploaded data (RFPs, proposals) is ever used for model training.

**Change Management & "Fear of Redundancy"**

Procurement teams, already skeptical of new tools, will resist adoption, fearing Aria is a threat to their job security. BCG research notes that frontline GenAI adoption has stalled at 51%.

**The "Liberation" Narrative**

The product's messaging and training are focused on up-skilling, not replacing. "Aria does the reading; you do the thinking." Aria automates the low-value grunt work to liberate humans to be more strategic.